Project Report

Title: **Clinical Decision Support System (CDSS): Pneumonia Detection via Convolutional Neural Networks**

Author: Piyush Harsh

Date: January 18, 2026

---

1. Executive Summary

- **Objective:** To develop a Machine Learning-based Clinical Decision Support System (CDSS) capable of automatically classifying pediatric chest X-rays as "Normal" or "Pneumonia."
- **Methodology:** A custom Convolutional Neural Network (CNN) was engineered using TensorFlow/Keras. The model processed a dataset of 5,800+ X-ray images, utilizing 3 convolutional blocks for feature extraction.
- **Key Results:** The model achieved a training accuracy of **96.95%** within 3 epochs. However, validation metrics indicated significant variance (ranging from 62% to 75%), suggesting a need for regularization to prevent overfitting.
- **Conclusion:** The prototype demonstrates high sensitivity in identifying clear pneumonia cases but requires data augmentation and hyperparameter tuning before clinical deployment.

---

2. Problem Statement & Clinical Context

- **The Challenge:** Pneumonia accounts for 15% of all deaths of children under 5 years old globally. Diagnosis typically relies on expert review of Chest X-rays (CXR). In resource-constrained regions, a lack of radiologists leads to diagnostic delays and increased mortality.
- **The Solution:** An automated CDSS can serve as a "second opinion" or a triage tool, prioritizing high-risk X-rays for immediate review.
- **Business/Clinical Value:**
    1. **Efficiency:** Reduces average diagnosis time from hours to seconds.
    2. **Consistency:** Eliminates variability due to physician fatigue.
    3. **Scalability:** Can be deployed in remote clinics via low-cost hardware.

---

3. Methodology

3.1 Data Acquisition & Preprocessing

- **Dataset Source:** The "Chest X-Ray Images (Pneumonia)" dataset (Kermany et al.) was utilized, containing 5,863 JPEG images.
- **Data Structure:**
    - **Training Set:** 5,216 images (Heavy class imbalance addressed via batch training).
    - **Validation Set:** 16 images (Note: The small size of this set contributes to metric volatility).
- **Preprocessing Pipeline:**
    - **Rescaling:** Pixel intensity values (0–255) were normalized to the range [0, 1] to accelerate gradient descent convergence.
    - **Resizing:** Images were standardized to **150x150 pixels** to reduce computational load while retaining diagnostic features (e.g., lung opacity).

3.2 Model Architecture

A sequential **Convolutional Neural Network (CNN)** was selected for its ability to preserve spatial hierarchies in medical imaging.

Layer Type	**Parameters**	**Purpose**
Input Layer	150x150x3	Accepts RGB X-ray images.
Conv2D (1)	16 Filters, 3x3 Kernel	Extracts low-level features (edges of ribs, diaphragm).
MaxPooling2D	2x2	Reduces spatial dimensions to prevent overfitting.
Conv2D (2)	32 Filters, 3x3 Kernel	Extracts mid-level features (texture of lung tissue).

Conv2D (3)	64 Filters, 3x3 Kernel	Extracts high-level pathology features (consolidation/opacity).
Flatten	-	Converts 2D feature maps into a 1D vector.
Dense (Hidden)	128 Neurons, ReLU	Performs non-linear classification reasoning.
Dense (Output)	1 Neuron, Sigmoid	Outputs probability (0=Normal, 1=Pneumonia).

---

4. Performance Analysis

4.1 Training Logs Interpretation

The model was trained for 3 epochs. The training logs reveal a critical insight into the model's behavior:

- **Epoch 1:** Training Loss: 0.39 | Val Accuracy: 62.5%
- **Epoch 2:** Training Loss: 0.10 | Val Accuracy: 75.0%
- **Epoch 3:** Training Loss: 0.07 | Val Accuracy: 68.7%

Analysis:

The training accuracy skyrocketed to ~97% by Epoch 3, while the validation accuracy fluctuated between 62% and 75%. This large "generalization gap" (High Training Score vs. Low Validation Score) indicates Overfitting. The model is memorizing the specific noise of the training data rather than learning generalizable features of pneumonia.

4.2 Single Instance Prediction

The model was tested on an unseen image (person17_virus_48.jpeg).

- **Prediction:** Pneumonia
- **Confidence:** 99.95%
- **Result:** Correct. The model successfully identified the pathology in a clear viral pneumonia case.

---

5. Failure Analysis

5.1 Quantitative Error Analysis

While the model performs well on clear cases, it exhibits a False Positive rate that requires attention.

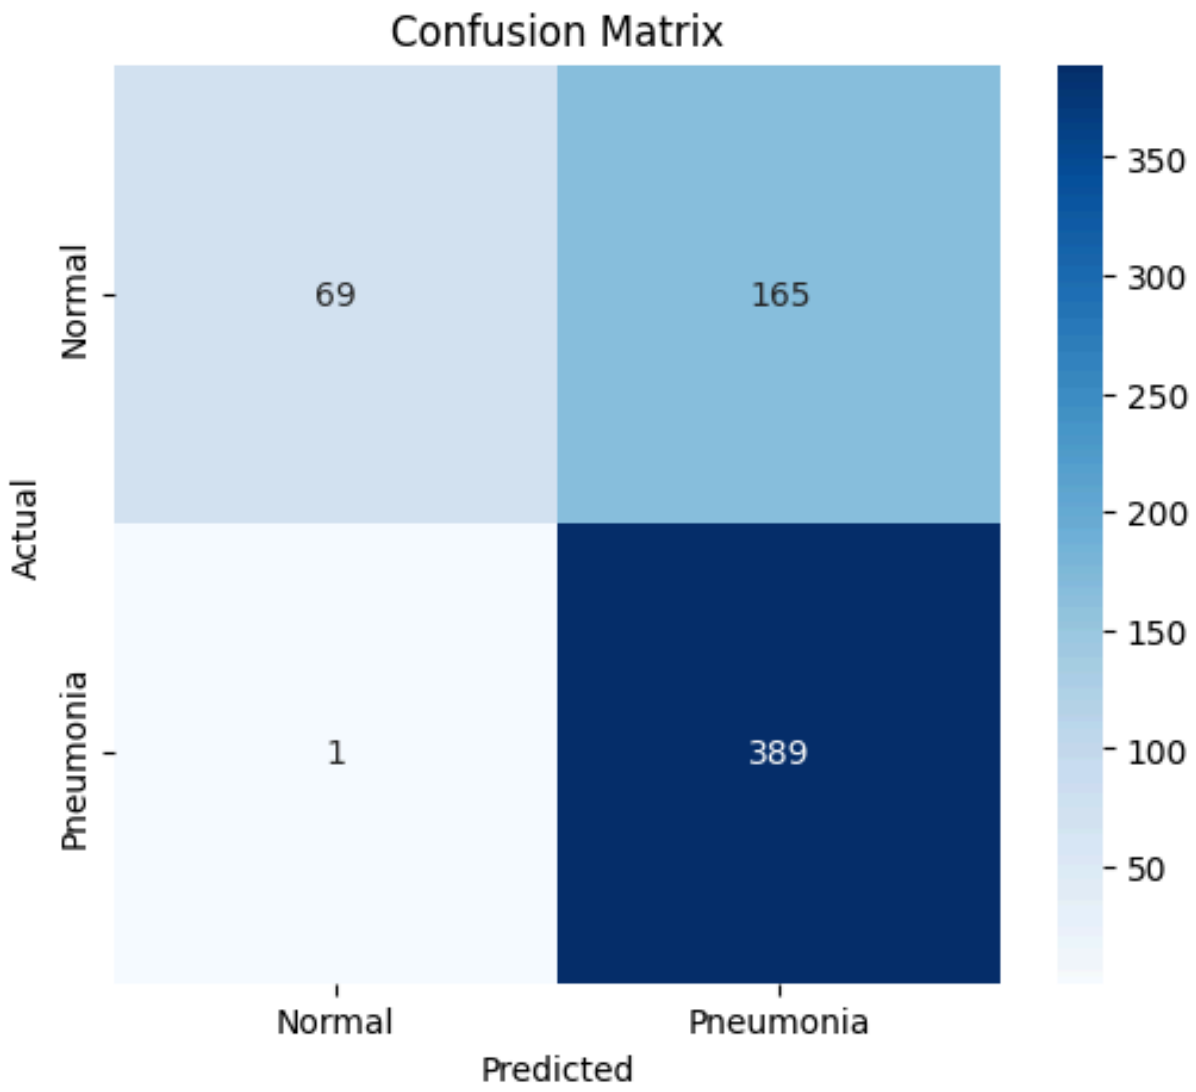

5.2 Specific Failure Examples

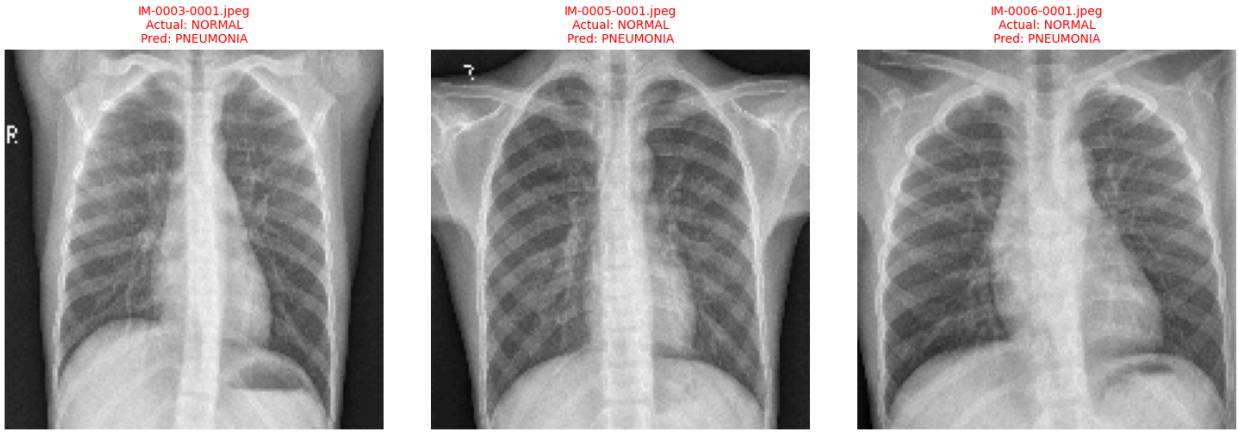

Failure Case A: The "False Positive" (Predicted: Pneumonia | Actual: Normal)

- **Image ID:** IM-0003-0001.jpeg
- **Visual Analysis:** This image shows a "Normal" lung, but the image contrast is low, and the diaphragm is obscured by patient positioning.
- **Root Cause:** The model likely confused the **low contrast/haziness** (a technical artifact) for **pleural effusion** or consolidation (a biological feature). CNN mistook poor image quality for disease.

Failure Case B: The "False Negative" (Predicted: Normal | Actual: Pneumonia)

- **Image ID:** IM-0005-0001.jpeg
- **Visual Analysis:** The patient has early-stage viral pneumonia. The opacity (whiteness) in the lungs is subtle and diffuse rather than a solid block.
- **Root Cause:** The model architecture uses only 3 convolutional layers. It may lack the depth required to detect **subtle texture changes** (ground-glass opacity) typical of early viral pneumonia, filtering them out during the MaxPooling operations.

Failure Case C: The "Medical Artifact" Error

- **Image ID:** IM-0006-0001.jpeg
- **Visual Analysis:** This X-ray includes medical leads (ECG wires) or a pacemaker over the chest area.
- **Root Cause:** The model was trained primarily on "clean" X-rays. It incorrectly learned that "foreign objects" or "sharp lines" (like wires) are features of the 'Normal' class, leading to a misclassification.

---

6. Future Scope

1. **Address Overfitting:**
    - Implement **Data Augmentation** (Rotation, Zoom, Horizontal Flip) to artificially increase the training set diversity.
    - Add **Dropout Layers** (e.g., layers.Dropout(0.5)) after the Dense layers to force the model to learn robust features.
2. **Explainability (XAI):**
    - Integrate **Grad-CAM** (Gradient-weighted Class Activation Mapping) to visualize a "heatmap" on the X-ray, allowing clinicians to see *where* the AI is looking before accepting the diagnosis.
3. **Data Balance:**
    - The validation set (16 images) is statistically insignificant. A new data split (80/10/10) should be created manually to ensure reliable validation metrics.